Beyond .*Script

Implementing A Language For The Web

Veit Heller

June 18, 2016

A thesis submitted for the degree of
B.Sc. of Applied Computer Science of
The University of Applied Sciences Berlin

For Meredith, Tobias and all the people who cope with me. Your undying support will not be forgotten.

Except where otherwise indicated, this thesis is my own original work.

Veit Heller

June 18, 2016

Abstract

The modern web is comprised of an abundance of very different beasts. Technologies that powered the first versions of the World Wide Web, such as HTML, CSS and JavaScript, and relatively new conceptions like TypeScript, CoffeScript, PureScript, ClojureScript, Elm, LASS, SCSS, Jade and Emscripten - to name but a few - are shaping the internet as we know it. There is one flaw that many of the new technologies have in common, as different as they may look and feel - they are mere preprocessors. In the end, it all boils down to the classic technologies again and we are left with the same limited capabilities we have had for the last twenty years.

This thesis presents a port of the zepto programming language to the web. It aims to work as seamlessly with existing technologies as possible.

Contents

Contents

Abbreviations

API Application Programming Interface. 8, 9

AST Abstract Syntax Tree. 2

DOM Document Object Model. 8, 11

FFI Foreign Function Interface. 8–11

GHC Glasgow Haskell Compiler. 2, 7

IR Intermediate Representation. 2

REPL Read-Eval-Print Loop. 7

W3C World Wide Web Consortium. 8

1 Introduction

> Controlling complexity is the
> essence of computer
> programming.
>
> *(*B. Kernighan*)*

1.1 Motivation

JavaScript has, since its inception, attracted a lot of controversy. This is rooted in various aspects of its design, from prototypal inheritance to operator precedence. Prototypal inheritance has the reputation of being counter-intuitive, though it is older than JavaScript, the first commonly known programming language that implements prototypal objects being Self.

* things get better
* es 6 and es7 thank god
* a lot of research funneled into it
* still a fundamental rethinking might be necessary

NOTE: This had to be moved from chapter 2 to match the current design of the thesis. Please excuse the lack of coherence to the part above.

1.1.1 Lisp

A common saying among programming language designers is that every programmer has written their own implementation of Lisp. There are a lot of different implementations of Lisp in the wild, even ones that compile to JavaScript[1].

The main reason for that is often cited to be the simplicity of the language on a parsing level. A simple Lisp can be implemented in less than one hundred lines of code, if no intermediate representation is generated. This is made possible by the unique property of Lisp of enclosing every statement in parentheses, where the first element

---

[1]such as ClojureScript, a backend of the Clojure compiler that targets JavaScript.

within those parentheses is the statement and the other elements are the arguments. It can be evaluated straight from a textual level, because things such as operator precedence and statement amiguity do not exist. In regular Lisp as specified in the initial paper by John McCarthy(McCarthy, 1960) only six special forms exist to allow not only for Turing-completeness, but also for expressiveness.

1.1.2 zepto

Zepto is a new Scheme implementation that aims to be as small as possible, to be able to target a lot of different backends. Currently, LLVM and Erlang Core[2] bindings are under development, the reference implementation is a simple interpreter that interprets code directly from the Abstract Syntax Tree (AST). This is slow but ensures a small interpreter size[3]. The compilers are written directly in zepto itself.

The small code base makes zepto a good target for porting it to the web. Further, ecause it is written in Haskell the code base was expected to be possibly almost entirely compilable to JavaScript using GHC-JS, a backend for the Glasgow Haskell Compiler (GHC) targetting JavaScript instead of native code. It offers many advanced features such as inlining of JavaScript into the code base using a technique called quasi-quoting, where a special character sequence delimits the inlined code, much like regular quotes. This tool set was expected to make the work of porting an existing language to the web as simple as possible.

Of course there are other reasons to use a functional language as an example. With both syntax and semantics differing wildly from JavaScript, this example makes way for languages more closely related to JavaScript also making their eventual way into the browser.

1.2 Goals of this Thesis

The primary goal of this thesis is to present a novel approach at implementing languages for the Web. This is exemplified by a sample implementation of a non-trivial functional programming language.

* functional because different

---

[2]Erlang Core is the Intermediate Representation (IR) of Erlang code before it is complied. Resources and documentation about it are sparse, it mostly seems to exist inside the BEAM's implementation.
[3]The entire codebase is only about 4000 lines of Haskell code.

* fairly different feature set

1.3 Structure of this Thesis

Chapter 2 examines related work in the field of cross-compilation into JavaScript and implementation of interpreters that are directly embeddable into larger systems. This includes desktop applications, game scripting engines and creative suites.

Chapter 3 gives an overview of the concept design and how the features are laid out to match the needs of both the goals of this thesis and the prototype itself.

Chapter 4 presents the system design and how the prototype integrates into existing web components.

Chapter 5 discusses the implementation, picking out different fundamental parts of the system and presents how they work.

Chapter 6 evaluates the prototype. This includes problems such as how well the integration of the system worked and how it compares to the reference implementation of zepto.

Chapter 7 gives a short summary of what was done and gives an outlook to what might happen with zepto, both the desktop and the JavaScript version, in the future.

2 Related Work

Practicality beats purity.

(T. Peters—The Zen of Python)

2.1 Existing Projects

2.2 Existing Standards

3 Concept Design

> Practicality beats purity.
>
> ─────────────────
>
> *(T. Peters—The Zen of Python)*

3.1 Construction Design

3.2 Additional Features

4 System Design

4.1 Integration into the Web Ecosystem

5 Implementation

> It always takes longer than you
> expect, even when you take into
> account Hofstadter's Law.
>
> *(*Hofstadter's Law*)*

The implementation philosophy philosophy of the port presented in this thesis has always been to reuse as much code from the reference implementation as possible. This guided the flow of design choises down a rather natural path and thus kept the implementation described here fairly short and relatively trivial.

5.1 Description of the Toolchain

The tooling uses GHCJS, which is a backend for the GHC compiler that targets JavaScript rather than native code (TODO: remove this from introduction). This makes cross-compiling the code base to JavaScript a rather simple undertaking.

* quasi-quoting
* jsbits
* as little dependence on it as possible

5.2 Description of the Implementation

As predicted in 1.1, the code base of zepto could be reused in almost its' entirety. What had to be rewritten was mostly related to the startup of the interpreter, because the regular paths into the code - either via a script being passed into it or launching an interactive Read-Eval-Print Loop (REPL)[1] - were unavailable in the browser context. Instead, a

---

[1] A REPL is an interactive code evaluation environment. Code is typed into a prompt and immediately evaluated. The convenience of such a short feedback loop is often used in the context of scripting languages and shells.

way of passing the sources from within script tags needed to be found. Further customizations include a Foreign Function Interface (FFI) to enable better cross-evaluation of JavaScript and the adaptation of existing Application Programming Interface (API)s, such as the Document Object Model (DOM).

5.2.1 The script tag

Initially, a DOM node walker was considered, but rejected relatively early because of two reasons: Firstly, it introduced a layer of complexity from within JavaScript code that would have likely made it brittle and hardly portable. Secondly, it would require a walk of the nodes every time a DOM element is inserted or replaced, which is a common occurence in modern interactive web applications.

As of November of 2015, the World Wide Web Consortium (W3C) specifies an API that simplifies this process for the programmer. Within their specification of the DOM4, the fourth specification of APIs for the Web, an object called MutationObserver is included which is able to register for DOM manipulations. Its main function will be triggered whenever a change occurs within the DOM part that it registered for listening to.

This simplifies the implementation of a listener to DOM events a great deal. Only minimal programming is required to configure the listener and to filter out all the nodes that are not script nodes of the type text/zepto[2].

A problem untended to with that method was nodes insert before the listener starts. This was resolved by singling out all the script tags that are present before the listener starts and applying the same filter/evaluation function to all of them. This also ensures that they are executed before any additional code (and possibly dependent) is passed into the zepto object.

The code was then included in the zepto singleton, which is the global interpreter object used for the management and interactivity of the zepto interpreter.

```
1  // the initial observer and the function it takes
2  zepto.observer = new MutationObserver(zepto.handleMutation);
3
4  // this function will get a list of mutations and apply handleDom to
       them
5  zepto.handleMutation = function(mutations) {
6    mutations.forEach(mutation => {
```

---

[2]This was chosen in analogy to the existing text/javascript node type

```
 7        mutation.addedNodes.map(zepto.handleDom);
 8      });
 9    }
10
11    // evaluate if it is a text/zepto node
12    zepto.handleDom = function(node) {
13      if (node.nodeName != "SCRIPT" || node.type != "text/zepto") {
14        return null;
15      }
16      return zepto.eval(node.innerHTML);
17    }
18
19    // execute this on startup
20    window.onload = () => {
21      let scripts = document.getElementsByTagName("script");
22      scripts.map(zepto.handleDom);
23    }
24    // the extra arguments signify recursive listening
25    zepto.observer.observe(document, {childList: true, subtree: true});
```

Listing 5.1: The final mutation observer code (simplified)

5.2.2 The FFI

The FFI is a central part of the port. If it weren't usable, none of the browser's capabilities could be used from within zepto, thus rendering the effort of bringing zepto into the browser effectively useless. The APIs of the Web are a big part of what it means to program for the browser, after all.

An initial sketch of the programming interface was extremely simplistic: a call to the function js could be called with a string as argument, representing the textual representation of the JavaScript program that should be run. It was piped to the JavaScript function eval and the function returned an affirmative truth value. Quasi-quoting larger blocks of JavaScript was also possible.

Of course this is unusable. The missing return value makes any effort of talking to an API impossible, as one could never yield any results. A different kind of return value is needed.

The obvious but most challenging to implement solution would be to infer a fitting

zepto type for every return value in JavaScript and return a result depending on that. While this could be seen as a rather elegant solution, it comes with its own set of caveats and exceptions, as the mapping between JavaScript and zepto values is not always obvious. A JavaScript object has too many properties that get lost in the process of translating it to zepto as to make it intuitive.

```
1  ; this would return an integer
2  (js "1 + 1")
3
4  ; this would return a hashmap
5  (js "{key: \"val\""})
6
7  ; this is problematic, because it will return an object
8  (js "new Error()")
```

Listing 5.2: The ideal FFI

Another problematic point is the implementation of JavaScript values in GHCJS. They are opaque datatypes, aliases for addresses and byte vectors. While zepto supports byte vectors and pointers, they are hardly a good representation for semantically rich prototypes as they only offer a glance into the underlying implementation of the JavaScript engine. While it is true that GHCJS itself provides methods for type coercion, they are crude and possibly error-prone.

A simpler method that is still mostly sensible came up: returning the string values of all of the values returned. While this places the burden of coercion into the programmer's hand, it also gives them the power to make their own decisions of how to deserialize values. Functions for deserializing the most common datatypes are included in the standard library of the JavaScript implementation of zepto, to aid the programmer in the process of finding the right methods of getting a value out of the FFI.

This still does not solve the problem of helping manage classes, but it empowers the programmer to find their own ways of serializing on the JavaScript side and deserializing on the zepto side to preserve the information they need in their specific programming context.

All of this needs an additional layer of abstraction to avoid unnecessary boiler plate, but it is stable enough for most purposes that zepto in JavaScript was used for yet.

```
1  ; the function string->number is a standard zepto function
2  (string->number (js "Math.pow(2, 32)"))
```

```
3
4  ; this is an example of how to resolve the earlier problem:
5  ; override the prototype of the object to return the value that is
      needed
6  (js "Error.prototype.toString = function() { return this.message; }")
7  (error (js "new Error(\"fatal error occured\")"))
```

Listing 5.3: The final form of the FFI

Implementing the JavaScript to zepto FFI was much simpler, as the interpreter is defined within the JavaScript environment. A call to the eval function of the zepto object with a string as argument will return in the execution of this piece of code and the return the textual representation of the zepto object so that the entire communication between the languages is string-based.

5.2.3 The DOM

After building the FFI, it was possible to implement the entire communication with the DOM in terms of calls to foreign functions and the parsing of their return values. This allows for a stable library, because it is unintrusive and does not interfer with existing JavaScript constructs.

Existing implementations often find it convenient to write a hybrid mix of JavaScript and zepto code that calls each other at certain points. This is, however, not advisable at the layer of libraries or utilities, because it is at risk of getting in the way of the job.

```
1  ; TODO: need to get this from work laptop
2  (module "dom"
3    (export
4      '("add-node" add-node)
5      '("get-elem" get-elem))
6
7    (define (add-node node)
8      ())
9
10   (define (get-elem elem)
11     ())))
```

Listing 5.4: A minimal version of a DOM module

6 Evaluation of the Prototype

> When I'm working on a problem,
> I never think about beauty. I
> think only how to solve the
> problem. But when I have
> finished, if the solution is not
> beautiful, I know it is wrong.
>
> *(*R. Buckminster Fuller*)*

6.1 Seamlessness of Integration

6.2 Test Against Standard Implementation of Zepto

7 Summary and Outlook

When I'm working on a problem,
I never think about beauty. I
think only how to solve the
problem. But when I have
finished, if the solution is not
beautiful, I know it is wrong.

(R. Buckminster Fuller)

* porting efforts
  * compiler efforts
  * classes

8 Conclusion

References

Aho, A. V., M. S. Lam, R. Sethi, and J. D. Ullman (2006). *Compilers: Principles, Techniques, and Tools*. Pearson Eudcation, Inc.

Fogus, M. (2013). *Functional JavaScript*. O'Reilly Media.

McCarthy, J. (1960). "Recursive Functions of Symbolic Expressions and Their Computation by Machine, Part I". In: *Communications Of The ACM* (cit. on p. 2).

Parr, T. (2010). *Language Implementation Patterns*. The Pragmatic Programmers, LLC.

Queinnec, C. (2003). *Lisp in Small Pieces*. Cambridge University Press.